

Network Service Interface Topology Service Distribution Mechanisms

Status of This Document

Group Working Draft (GWD), Informational (I).

Copyright Notice

Copyright © Open Grid Forum (2008-2013). Some Rights Reserved. Distribution is unlimited.

Abstract

This document describes a normative schema which allows the description of service plane objects required for the Network Service Interface Connection Service. Additionally it describes a set of distribution mechanisms for the network topology descriptions.

Contents

Abstract	1
Contents	1
1 Representation of Network Topologies	3
1.1 Introduction to concept of STPs	3
1.2 Identifying STPs in a request	3
1.3 Explicit routing using STPs	4
1.4 Requirements for Topology Descriptions	4
1.5 Topology Identifiers	5
2 NSI Terminology	6
3 Distribution of NSI Topology	6
3.1 Transport and Service plane relationship	6
3.2 Elements of a Topology Exchange Mechanism	6
3.3 Topology Exchange Implementations	7
	1

3.4 Summarizing Topology Distribution	8
4 Contributors	9
5 Acknowledgments	9
6 Intellectual Property Statement	10
7 Disclaimer	10
8 Full Copyright Notice	10
References	12
Normative References	12
Informative References	12

1 Representation of Network Topologies

In order to use the NSI Connection Service, some form of topology representation is required. In this section we explain the topology representation requirements for the NSI Connection Service. A diagram that provides some generic insight into NSI Topology is provided in figure 1.

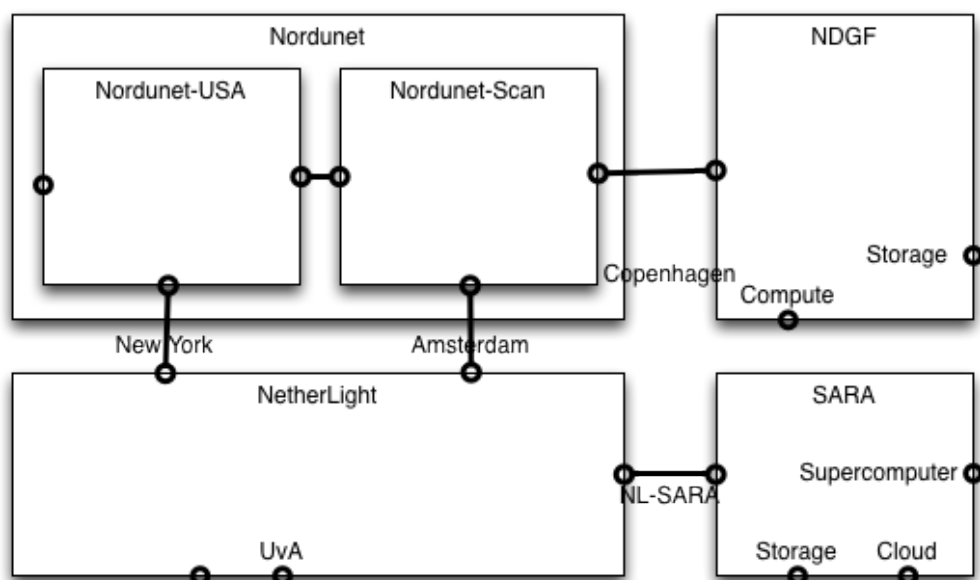


Figure 1: Abstract view of an example NSI Topology

1.1 Introduction to concept of STPs

The basic network topology for the Network Service Interface (NSI) consists of networks, points and connections. The NSI can be used to request a connection between two different points, which is then implemented using the connection(s) between those points. Since each of the points in the network topology can terminate a network service, they are called Service Termination Points (STPs). The figure above contains several of these, for example the Storage point at SARA, another example is the connection point between the network on the edge of the SARA network which connects to the Netherlight network.

1.2 Identifying STPs in a request

The STPs in the Network Topology generally have two different roles:

- Endpoints within the network – points which are of interest to users, used as source and destination points
- One Part of a Service Demarcation Point (SDP) – meaning connections to other networks, used as transit points

For the NSI network service, knowledge of the connections between networks is necessary to enable pathfinding. It is not strictly necessary to know all of the user endpoints in a network, as long as it is known in which network the endpoint is. It is assumed that once the network is reached, the endpoint within the network can be reached as well.

To allow for pathfinding it is necessary that the Topology Service can identify the Networks, and the connections between those Networks, the SDPs. It is not necessary to distribute all of the endpoints within the network, which makes the distribution process much simpler.

This means that an NSI Connection Service reservation request must contain the source and destination network, as well as the connection points within those respective networks.

1.3 Explicit routing using STPs

In a regular request only the source and destination STPs and their networks are specified. The selected path between those STPs is left to the NSA managing the request. In NSI v2.0 it is also possible to steer the requested path into a specific direction by defining intermediate STPs that the path must touch.

Instead of a normal request with just a source and destination STP, the explicitly routed request will contain a path element which contains an indication of the path that should be taken through the NSI Network.

If the Path object contains a description of the complete path end-to-end then this is simply a question of availability. However, problems can arise if the path object only contains a single “explicit route object”(ERO) that it must touch. With the current NSI implementation (v1) and its bi-directional model, there is no way to know from which way to cross that ERO. The proposal is to use uni-directional path elements, to avoid ambiguity in the path direction (the uni-directionality of the ERO implicitly defines the direction).

1.4 Requirements for Topology Descriptions

Taking the description above into account, and with the general idea of the Network Service Interface in mind, we come to the following requirements for the network topology description:

Scalable : An NSA does not need to be aware of all STPs in other networks;

Compact : A Topology description should be able to group individual data transport capabilities in one object rather than specifying each possible VLAN for example;

Abstract : The topology description should list the connections between domains, not how these connections are implemented;

Compatible : It should be possible to relate the NSI topology to other topologies, e.g. as used for monitoring. Preferably the NSI topology should be compatible with the NML topology;

Flexible : The topology description should support future extensions, e.g. different connection types (unidirectional, bidirectional, or multipoint connections), different levels of abstraction of the network (subtopologies), or multihoming scenarios;

Unambiguous : The direction of an STP in an ERO should be unambiguous.

1.5 Topology Identifiers

To satisfy the unambiguity and flexibility requirements, we propose to describe *STPs* as two unidirectional *Ports*, since these have a single direction, there can be no misunderstanding. These unidirectional ports also easily allow point-to-multipoint requests.

To satisfy the compact requirement, we propose to allow *PortGroups* over *Ports*. A *PortGroup* groups together several *Ports* which have a single identifying attribute, for example a VLAN label.

To satisfy the scalability and flexibility requirement we propose to add the Topology identifier as an added context for an STP in a request. This makes most of the global path through the network clear, without having specific knowledge about the internal endpoints. These can be handled by the NSAs responsible for those Topologies.

To satisfy the compatibility and flexibility requirements we propose to use full URIs for each component in a request. Globally unique identifiers make it possible to have delegated subtopologies without having to rewrite identifiers for example.

This means that a connection request for NSI SHOULD use the following tuple to identify a source or destination STP:

(Topology ID, source PortGroup ID, sink PortGroup ID)

2 NSI Terminology

The NSI Connection Service requires topology descriptions to do pathfinding. In order to do that some representation of the topology is required. Once represented, some form of topology distribution is also needed. This document describes some requirements for the NSI Topology Service, suggests a short-term implementation and a strategy for better long-term support.

In the first section we describe what is necessary for the topology to support, what kind of elements should be in there. In the next section we describe the distribution requirements, some possible solutions and a recommended solution for the short-term and also for the longer term

3 Distribution of NSI Topology

Some form of Topology distribution is required in order for an inter-domain NSI network to function. In NSI 1.0 this process was performed out-of-band, mostly through e-mail. For NSI 2.0 we take the opportunity to define an NSI Topology Service for NSI topology exchange, which can support the NSI Connection Service.

3.1 Transport and Service plane relationship

The NSI Connection Service is implemented on Network Service Agents (NSAs), which together form a network and tree-like structure. This graph represents how reservation requests would propagate through the network, but not necessarily reflects the transport-plane. One NSA may be an aggregation point for other NSAs, not visible from the outside.

The messaging between the NSAs will happen on the service plane, which is separated from the transport plane.

3.2 Elements of a Topology Exchange Mechanism

There are three main elements of topology exchange:

Bootstrapping Topology Exchange To start the initial Network Service the NSAs must be able to find each other, in order to communicate details about the network. So some form of bootstrapping is required, with initial synchronization between domains on both the service plane and the transport plane, i.e. the NSAs of both domains must be able to contact each other, and the details of transport plane connections between them have to be synchronized as well.

Initiating a transport plane connection between two networks is not a frequent occurrence, and a longer process, involving out-of-band (for NSI) contact. Part of that process can be that the networks also communicate the access details for the NSAs, thus forming an NSA relationship.

Expanding the Topology Exchange Once the neighboring (on the control-plane) NSAs have exchanged details, they can also distribute details about the rest of the network, both the control plane details and connectivity, but also some transport details.

Update Mechanism The transport network is not static, and links are added or removed from time to time. An update mechanism is thus required to inform other NSAs about these kinds of events.

3.3 Topology Exchange Implementations

The above mechanisms can be implemented in five different ways:

Centralized Manual Distribution An initial attempt at topology distribution in the Automated GOLE demonstration was through a central maintainer. This maintainer collected all topology information from the networks involved, gathered all the topology data and sent out a topology file through e-mail. The network maintainers would then download the attachment and insert it into their provisioning system. Updates to the topology were all handled through the central maintainer, distributed through e-mail.

While this system worked initially, it soon ran into scaling problems. This system also does not allow to have a good way of doing automatic updates or insertion.

Version Controlled Distribution The Automated GOLE demonstration has transitioned into a different distribution mechanism using a Git source code repository, available on GitHub. This mechanism also still has a central maintainer, this also allows networks to manipulate their own topology information. The distribution mechanism is either directly from the GitHub website, or through the git version control system itself. The git system has the added advantage that it is a distributed version control system, so it is not required to download directly from GitHub.

Bootstrapping and updating all happens through the git system itself.

PerfSonar Lookup Service The PerfSonar monitoring system suite also contains a service for looking up information. This service uses a “home Lookup Service”(hLS) where metadata of information is registered, which is then uploaded into the “global Lookup Service”(gLS) (this can be cloud of services).

The retrieval of information happens by first querying the gLS, then the relevant hLS, followed by the service where the actual information is stored.

Since topology information would be stored locally, no update mechanism is necessary, except for location changes of the services itself. However, this method of storing and lookup does require full connectivity between all NSAs to provide and retrieve information, which may not be possible.

A soon to be released (target Dec 2012) updated PS Lookup service will incorporate subscription capabilities, which allow an hLS to push information to a remote hLS and have it cached locally there. By adapting the new Lookup service to store topology, and selectively managing the subscription of data, full connectivity between NSAs would not be necessary to disseminate global topology.

HTTP Distribution A common way of distributing information is using the HTTP protocol. The topology files would include links to topology description files, which would allow other domains to directly fetch the topology description.

However, as with the Lookup Service, this requires direct access between NSAs which may not be possible.

A Peer-to-Peer Distribution Protocol Another method is to define a new protocol for NSI topology distribution. As explained above the protocol would only require a small set of primitives, and would work directly between peering NSAs.

Once an NSA comes up, it contacts its neighbors to request the topology information that they know about, and subscribes to future topology updates. These updates are propagated in a peer-to-peer fashion through the whole network.

The end-result would be that all the NSAs have a global view of the topology, with only local interaction.

3.4 Summarizing Topology Distribution

Above we have described five different topology distribution mechanisms, from these five only the version controlled and peer-to-peer systems fit the requirements that the NSI has, others require too much manual operation, or require globally reachable NSAs. The peer-to-peer system would be the best solution with minimal interaction between systems, and no reliance on outside mechanisms. For practical reasons we believe the best solution right now is a version controlled distribution system, and in time evolve to a peer-to-peer distribution protocol.

4 Contributors

Jeroen J. van der Ham (Editor)

Faculty of Science, Informatics Institute, University of Amsterdam

Science Park 904, 1098 XH Amsterdam

The Netherlands

Email: vdham@uva.nl

5 Acknowledgments

The author would like to thank the OGF NSI, NML and GLIF DTOX groups for their input.

6 Intellectual Property Statement

The OGF takes no position regarding the validity or scope of any intellectual property or other rights that might be claimed to pertain to the implementation or use of the technology described in this document or the extent to which any license under such rights might or might not be available; neither does it represent that it has made any effort to identify any such rights. Copies of claims of rights made available for publication and any assurances of licenses to be made available, or the result of an attempt made to obtain a general license or permission for the use of such proprietary rights by implementers or users of this specification can be obtained from the OGF Secretariat.

The OGF invites any interested party to bring to its attention any copyrights, patents or patent applications, or other proprietary rights which may cover technology that may be required to practice this recommendation. Please address the information to the OGF Executive Director.

7 Disclaimer

This document and the information contained herein is provided on an “As Is” basis and the OGF disclaims all warranties, express or implied, including but not limited to any warranty that the use of the information herein will not infringe any rights or any implied warranties of merchantability or fitness for a particular purpose.

8 Full Copyright Notice

Copyright © Open Grid Forum (2008-2013). Some Rights Reserved.

This document and translations of it may be copied and furnished to others, and derivative works that comment on or otherwise explain it or assist in its implementation may be prepared, copied, published and distributed, in whole or in part, without restriction of any kind, provided that the above copyright notice and this paragraph are included as references to the derived portions on all such copies and derivative works. The published OGF document from which such works are derived, however, may not be modified in any way, such as by removing the copyright notice or references to the OGF or other organizations, except as needed for the purpose of developing new or updated OGF documents in conformance with the procedures defined in the OGF Document Process, or as required to translate it into languages other than English. OGF, with the approval of its board, may remove this restriction for inclusion of OGF document content for the purpose of producing standards in cooperation with other international standards bodies.

The limited permissions granted above are perpetual and will not be revoked by the OGF or its successors or assignees.

References

Normative References

- [NSI-FRAMEWORK] Guy Roberts, Tomohiro Kudoh, Inder Monga, Jerry Sobieski, and John Vollbrecht. Network Services Framework v1.0 GFD.173 (Informational), December 2010. URL <http://ogf.org/documents/GFD.173.pdf>.
- [URN-OGF-NETWORK] Freek Dijkstra, and Jeroen van der Ham. A URN Namespace for Network Resources. GWD-I *draft-gwdi-urn-ogf-network* (Work in Progress), September 2012. URL <https://forge.ogf.org/sf/go/doc16260>.
- [ISO 8601] Data elements and interchange formats – Information interchange – Representation of dates and times. ISO 8601:2004 (Third edition), December 2004. Section 4.3.2 (a), Complete representations of a date and time. Calendar date in basic format. URL http://www.iso.org/iso/home/store/catalogue_ics/catalogue_detail_ics.htm?csnumber=40874.
- [NML] Jeroen van der Ham, Freek Dijkstra, Roman Łapacz, and Jason Zurawski. Network Markup Language Base Schema version 1. GWD-R-P *draft-gwdrp-nml-base* (Work in Progress), January 2013. URL <https://redmine.ogf.org/attachments/46/nml-base.pdf>.
- [RDF-XML] Dave Beckett (editor) RDF/XML Syntax Specification (Revised) W3C Recommendation 10 February 2004. URL <http://www.w3.org/TR/rdf-syntax-grammar/>.
- [RFC 2119] Scott Bradner. Key words for use in RFCs to Indicate Requirement Levels. RFC 2119 (Best Current Practice), March 1997. URL <http://tools.ietf.org/html/rfc2119>.
- [RFC 3492] A. Costello Punycode: A Bootstring encoding of Unicode for Internationalized Domain Names in Applications (IDNA) RFC 3492 (Standards Track), March 2003 URL <http://tools.ietf.org/html/rfc3492>.
- [RFC 3986] Tim Berners-Lee, Roy T. Fielding, and Larry Masinter. Uniform Resource Identifier (URI): Generic Syntax RFC 3986 (Standards Track), January 2005. URL <http://tools.ietf.org/html/rfc3986>.
- [XML] Henry S. Thompson, David Beech, Murray Maloney and Noah Mendelsohn XML Schema Part 1: Structures Second Edition W3C Recommendation 28 October 2004. URL <http://www.w3.org/TR/xmlschema-1/>.

Informative References

- [Dijkstra13] Freek Dijkstra, et al. Experimental Features for NML 1. Work in Progress.
- [GFD.165] Paola Grosso, Aaron Brown, Aurélien Cedeyn, Freek Dijkstra, Jeroen van der Ham, Anand Patil, Pascale Primet, Martin Swany, and Jason Zurawski. Network Topology Descriptions in Hybrid Networks GFD 165 (Informational), March 2010. URL <http://www.ogf.org/documents/GFD.165.pdf>.
- [RFC 6350] Simon Perreault. vCard Format Specification RFC 6350 (Standards Track), August 2011. URL <http://tools.ietf.org/html/rfc6350>.
- [RFC 6351] S. Perreault. xCard: vCard XML Representation RFC 6351 (Standards Track), August 2011. URL <http://tools.ietf.org/html/rfc6351>.
- [RDFVCARD] Harry Halpin, Renato Iannella, Brian Suda, Norman Walsh Representing vCard Objects in RDF W3C Member Submission 20 January 2010. URL <http://www.w3.org/TR/vcard-rdf/>.